



Secure VM-Based Agent Labs on Apple Silicon with UTM, Ubuntu ARM64, and Rootless Containers

Executive summary

A reproducible, security-conscious way to experiment with autonomous “do-things-on-your-machine” agents on an Apple Silicon Mac is to treat the agent runtime as **untrusted** and run it inside a **Linux VM** with **NAT-based networking, no host file sharing**, and a **rootless container runtime** for further isolation. UTM provides the VM layer (and can be automated via Vagrant + the UTM provider), while projects in the OpenClaw ecosystem provide multiple “in-VM” deployment patterns: (a) Docker/Compose flows that can be adapted to loopback-only exposure, (b) explicit **rootless Podman** scripts with systemd Quadlet support, and (c) hardening playbooks that add **egress allowlisting** (Squid), reverse-proxy TLS, and non-root execution.

1

The best “building blocks” (closest to your preferred stack: **UTM + Ubuntu Server ARM64 → rootless Docker → OpenClaw**) are:

- **UTM + Vagrant automation (ARM64 base box)**: a concrete, reproducible UTM-on-Apple-Silicon workflow using `vagrant_utm`, including an `utm/ubuntu-24.04` base box marked **arm64**, with example `Vagrantfile` configuration and screenshots of UTM console setup. 2
- **Rootless Docker provisioning recipe you can transplant into your UTM VM**: `lima-vm/lima`'s `templates/docker.yaml` is an explicit, working rootless Docker installation script (disable rootful Docker service, install `uidmap` / `dbus-user-session`, run `dockerd-rootless-setuptool.sh`, verify `rootlesskit`). You can adapt the same steps into cloud-init, a Vagrant shell provisioner, or an Ansible task set inside your Ubuntu VM. 3
- **Defense-in-depth hardening for the agent runtime**: a Tier-3-style hardened deployment that runs the agent in **rootless Podman**, adds **Squid egress allowlists**, and uses a reverse proxy for HTTPS—useful if you want to aggressively limit outbound traffic and reduce credential exposure. 4

What follows is a curated set of public repos and guides (6-12) that most directly support a secure, VM-first workflow, plus practical adaptation notes for your specific UTM+Ubuntu ARM64 target.

Reference lab architecture and hardening concepts

A “VM-first, container-second” lab separates risk into layers:

- **VM boundary (UTM)**: limits direct access to macOS host state and keeps agent execution away from host credentials and files by default. UTM’s Apple networking “shared” mode is explicitly documented as **NAT-based sharing with the host** in its scripting reference. 5
- **Network posture**: UTM documents two primary network modes—“Shared Network” and “Bridged”—and calls bridged “for advanced users.” Shared mode is recommended for new VMs; from

a lab perspective, the key property is to avoid L2 adjacency to the building network unless you intentionally need it. ⁶

- **Rootless containerization inside the VM:** reduces the blast radius if the agent breaks out of its own process sandbox. `lima-vm/lima`'s Docker template explicitly installs Docker and then configures **rootless** operation and validates that `rootlesskit` is running. ⁷
- **Outbound (egress) control:** for agents, outbound calls are often the most dangerous (data exfiltration + prompt injection supply chains). A hardened playbook pattern is to force agent traffic through an allowlisting proxy like Squid. ⁴
- **Expose services safely:** OpenClaw's own docs emphasize keeping the Gateway bound to loopback when using tailnet-based exposure mechanisms, and provide guidance for Tailscale Serve/Funnel behavior and auth requirements. ⁸

For UTM operationally, headless operation is supported by removing display devices and controlling the VM via UI or scripting. This matters because many security-first labs prefer **no GUI integrations** and minimal host→guest convenience features. ⁹

High-quality repositories and public projects

The table below includes repos/projects and a small number of community guides. "URL" is provided as a clickable citation link (instead of a raw URL). Licenses are recorded only where visible in sources reviewed; otherwise marked "Not determined."

Repo / project	URL	Primary author/org	License	Target arch (ARM/x86)	VM tech used	Provisioning tools
openclaw/openclaw	Link ⁸	Peter Steinberger ¹⁰	MIT ¹⁰	Not specified (multi-platform positioning) ⁸	None (runtime + Docker/Podman artifacts) ⁸	Docker Compose + scripts; includes rootless Podman setup script ¹¹
openclaw/openclaw-ansible	Link ¹⁵	openclaw.org ¹⁵	MIT ¹⁵	Debian/Ubuntu Linux (host target) ¹⁵	None (intended to run on Linux; VM recommended for macOS users) ¹⁵	Ansible playbook + installer scripts ¹⁵

Repo / project	URL	Primary author/org	License	Target arch (ARM/x86)	VM tech used	Provisioning tools
Next-Kick/openclaw-hardened-ansible	Link 4	Next-Kick 4	Not determined (README says "provided as-is") 4	Arch + Debian/Ubuntu targets 4	None (runs on a target machine—ideal inside a VM) 4	Ansible + scripts (deploy.sh) 4
naveenrajm7/vagrant_utm	Link 16	Naveen Raj M 16	MIT 16	UTM-based boxes can be ARM or x86 (docs mention x86 possible via x86 boxes; UTM supports multiple architectures) 17	UTM (Vagrant provider) 16	Vagrant provider plugin; Vagrant provisioning supported 18
SrivatsaRv/vagrant-utm-demo	Link 20	SrivatsaRv 20	CC0-1.0 20	Not specified 20	UTM + Vagrant 20	Vagrantfile + provision.sh + Docker Compose 20
naveenrajm7/K8s-utm	Link 21	Naveen Raj M 21	Not determined (not shown in source excerpt) 21	Not specified 21	UTM + Vagrant (3-node cluster) 21	Vagrant + Ansible playbooks 21
andreesg/openclaw-terraform-hetzner	Link 22	andreesg 22	MIT 22	Cloud VPS (not tied to Apple Silicon; arch depends on chosen server type) 22	Cloud VM (Hetzner VPS) 22	Terraform + cloud-init templates + deployment scripts 22

Repo / project	URL	Primary author/org	License	Target arch (ARM/x86)	VM tech used	Provisioning tools
dazeb/proxmox-openclaw-installer	Link <small>23</small>	dazeb <small>23</small>	Not determined (not shown) <small>23</small>	Depends on Proxmox host + chosen cloud image (Ubuntu 24.04 cloud image referenced) <small>23</small>	Proxmox VE (automated VM provisioning) <small>23</small>	Shell installer + Proxmox cloud-init/snippets workflows <small>23</small>
lima-vm/lima (plus its Docker template)	Link <small>24</small>	lima-vm <small>24</small>	Apache-2.0 <small>24</small>	Explicitly supports Apple Silicon workflows; Docker template includes Rosetta notes for running <code>linux/amd64</code> containers on Apple Silicon <small>3</small>	QEMU / Apple Virtualization.framework (via Lima) <small>24</small>	YAML templates + cloud-init style provisioning scripts <small>3</small>
Small Sharp Software Tools: UTM + Vagrant tutorial (guide)	Link <small>25</small>	Brian P. Hogan (site author) <small>25</small>	N/A (tutorial content; not a repo) <small>25</small>	Uses <code>utm/ubuntu-24.04</code> base box marked arm64 <small>25</small>	UTM + Vagrant + <code>utmctl</code> <small>25</small>	Vagrantfile + optional shell provisioning snippet (example installs Tailscale) <small>25</small>

Top three blueprints adapted to UTM + Ubuntu ARM64 VM → rootless Docker → OpenClaw

UTM reproducibility via Vagrant + `vagrant_utm` and an ARM64 Ubuntu base box

This is the most direct “UTM-first and reproducible” path: install Vagrant + the UTM provider plugin and spin up an Ubuntu VM from the `utm/ubuntu-24.04` base box, which the tutorial’s output shows as `provider: utm (arm64)`. 2

A minimal starting point, derived from the tutorial’s `Vagrantfile`, looks like this (modify CPU/RAM/disk as needed): 25

```

Vagrant.configure("2") do |vagrant|
  hostname = "agentlab-ubuntu"

  vagrant.vm.define hostname, primary: true do |machine|
    machine.vm.box = "utm/ubuntu-24.04"
    machine.vm.hostname = hostname

    machine.vm.provider "utm" do |u|
      u.name = hostname
      u.cpus = 2
      u.memory = 4096
      # The tutorial demonstrates directory_share_mode = "virtFS"
      # If you want "no shared folders", leave sharing off rather than enabling
      # it.
      # u.directory_share_mode = "virtFS"
    end

    # Add provisioning here (rootless Docker + OpenClaw).
  end
end

```

Key pitfalls to anticipate:

- The tutorial maps SSH for access (`vagrant ssh`) and explains that the VM can be headless by default; if you want an on-demand console, it shows how to add a `virtio-gui-pci` display in UTM and includes screenshots. ²⁵
- The same tutorial demonstrates enabling directory sharing via `directory_share_mode = "virtFS"`; for an agent sandbox, that “convenience” can become the main exfil path, so treat it as an opt-in. ²⁵

Rootless Docker inside the VM by transplanting Lima's proven install steps

If you want rootless Docker specifically (not just “Docker inside a VM”), the most explicit and reusable reference implementation I found is Lima's `templates/docker.yaml`. It installs Docker, disables the rootful Docker system service, installs `uidmap` and `dbus-user-session`, then installs rootless Docker via `dockerd-rootless-setuptool.sh` and switches the Docker context to `rootless`. It also probes for `rootlesskit` to confirm that the rootless runtime is active. ⁷

You can translate its core steps into a Vagrant “shell” provisioner, cloud-init, or Ansible task. The sequence below is directly aligned with Lima's template logic: ⁷

```

# inside Ubuntu VM
curl -fsSL https://get.docker.com | sh

# Prefer rootless; Lima explicitly disables the system docker service

```

```

sudo systemctl disable --now docker

# Dependencies used by Lima
sudo apt-get update
sudo apt-get install -y uidmap dbus-user-session

# Rootless install (runs as your user)
systemctl --user start dbus
dockerd-rootless-setuptool.sh install
docker context use rootless

```

If you later decide you need multi-arch container runs (e.g., `linux/amd64` containers on Apple Silicon), Lima's template includes explicit notes about Rosetta-enabled VMs and running `--platform=linux/amd64` with Rosetta caching. That's a useful "future knob" for cross-arch experimentation even if you start pure ARM64. ⁷

Hardening with egress allowlisting + rootless containers using the Tier-3 Ansible playbook

If you want your lab to resist the common failure mode ("agent got tricked into curling random scripts / exfiltrating secrets"), the most concrete public project here is `Next-Kick/openclaw-hardened-ansible`. It explicitly deploys using **rootless Podman**, and adds **Squid proxy egress allowlisting**, plus a reverse proxy (Caddy) and additional monitoring and identity automation. ⁴

Even if you keep "rootless Docker preferred" as your default, this repo is valuable as a reference for what "serious" agent hardening looks like:

- Rootless containers + non-privileged user
- Domain allowlists for outbound HTTP(S)
- TLS termination and controlled UI access
- "Agent never sees real API keys" via a credential broker pattern (LiteLLM) ⁴

From an adaptation standpoint: run this playbook **inside** your UTM Ubuntu VM to keep macOS isolated, then treat the VM as disposable (snapshots + rebuild). ²⁶

UTM-specific notes for isolation, networking, and snapshots

Networking defaults and what "shared" really means

UTM's Apple "Network Mode" docs describe Shared Network as host-routed traffic where guest and host can see each other without extra config, and "Bridged" as for advanced users. ⁶

UTM's scripting reference is even more explicit: `shared` is "NAT based sharing with the host," while `bridged` is "Bridged to a host interface." ⁵

For a secure agent lab, this makes "shared/NAT" a natural baseline: you get outbound internet access without placing the VM directly onto the building's Wi-Fi broadcast domain. ²⁷

If you need to selectively expose services from the VM to the host, UTM documents port forwarding for QEMU backend VMs under specific network modes. ²⁸

(Separately, the Small Sharp tutorial relies on Vagrant's SSH port mapping to access the headless Ubuntu VM via `vagrant ssh`.) ²⁵

Snapshots as a safety tool

At the disk-image level, QEMU's own documentation notes that VM snapshots require a writable `qcow2` disk image. ²⁹

If your UTM workflow uses `qcow2` disks (common for QEMU-backed VMs), you can leverage snapshot semantics to "rewind" after risky experiments. ²⁹

Where UTM's UI surfaces are limited, there is at least one public companion project: [Metamogul/UTM-Snapshot-Manager](#), described as a snapshot-management companion app that relies on QEMU snapshot management for `.qcow` formatted disks. ³⁰

Operational mode: headless and minimal integration

UTM supports headless mode by removing display devices and controlling the VM via UI or scripting; it recommends keeping at least one serial device available for VM communication. ⁹

This aligns with a "secure lab" posture: fewer host-integrated features, fewer bidirectional surfaces. ³¹

Agent-specific risks and mitigations reflected in the projects

OpenClaw's own repo describes a security model in which non-main sessions (e.g., group/channel contexts) can run inside per-session Docker sandboxes, and it provides guidance for not exposing the Gateway directly and for using tailnet-only exposure patterns (Serve/Funnel) that keep the Gateway bound to loopback with auth constraints. ⁸

Automation installers in the ecosystem increasingly treat virtualization and hardening as first-class. [openclaw/openclaw-ansible](#) explicitly disables bare-metal macOS support due to security risk, recommends against bypassing that check, and highlights a firewall-first posture (UFW + Docker isolation), fail2ban, unattended upgrades, non-root execution, and restricted public ports (SSH and Tailscale). ¹⁵

Finally, the "Tier 3" hardening playbook demonstrates the direction the community is moving for safer experimentation: rootless containers, outbound allowlists, reverse-proxy gating, and explicit device-approval workflows. ⁴

¹ ⁴ ²⁶ <https://github.com/Next-Kick/openclaw-hardened-ansible>
<https://github.com/Next-Kick/openclaw-hardened-ansible>

² ²⁵ <https://smallsharpsoftwaretools.com/tutorials/utm-vagrant/>
<https://smallsharpsoftwaretools.com/tutorials/utm-vagrant/>

³ ⁷ <https://raw.githubusercontent.com/lima-vm/lima/master/templates/docker.yaml>
<https://raw.githubusercontent.com/lima-vm/lima/master/templates/docker.yaml>

- 5 27 <https://docs.getutm.app/scripting/reference/>
https://docs.getutm.app/scripting/reference/
- 6 <https://docs.getutm.app/settings-apple/devices/network/>
https://docs.getutm.app/settings-apple/devices/network/
- 8 <https://github.com/openclaw/openclaw>
https://github.com/openclaw/openclaw
- 9 31 <https://docs.getutm.app/advanced/headless/>
https://docs.getutm.app/advanced/headless/
- 10 <https://raw.githubusercontent.com/openclaw/openclaw/main/LICENSE>
https://raw.githubusercontent.com/openclaw/openclaw/main/LICENSE
- 11 13 <https://raw.githubusercontent.com/openclaw/openclaw/main/docker-compose.yml>
https://raw.githubusercontent.com/openclaw/openclaw/main/docker-compose.yml
- 12 <https://raw.githubusercontent.com/openclaw/openclaw/main/setup-podman.sh>
https://raw.githubusercontent.com/openclaw/openclaw/main/setup-podman.sh
- 14 <https://raw.githubusercontent.com/openclaw/openclaw/main/docker-setup.sh>
https://raw.githubusercontent.com/openclaw/openclaw/main/docker-setup.sh
- 15 <https://github.com/openclaw/openclaw-ansible>
https://github.com/openclaw/openclaw-ansible
- 16 18 19 https://github.com/naveenrajm7/vagrant_utm
https://github.com/naveenrajm7/vagrant_utm
- 17 https://github.com/naveenrajm7/vagrant_utm/discussions/15
https://github.com/naveenrajm7/vagrant_utm/discussions/15
- 20 <https://github.com/SrivatsaRv/vagrant-utm-demo>
https://github.com/SrivatsaRv/vagrant-utm-demo
- 21 <https://github.com/naveenrajm7/K8s-utm>
https://github.com/naveenrajm7/K8s-utm
- 22 <https://github.com/andreesg/openclaw-terraform-hetzner>
https://github.com/andreesg/openclaw-terraform-hetzner
- 23 <https://github.com/dazeb/proxmox-openclaw-installer>
https://github.com/dazeb/proxmox-openclaw-installer
- 24 <https://github.com/lima-vm/lima>
https://github.com/lima-vm/lima
- 28 <https://docs.getutm.app/settings-qemu/devices/network/port-forwarding/>
https://docs.getutm.app/settings-qemu/devices/network/port-forwarding/
- 29 <https://qemu-project.gitlab.io/qemu/system/images.html>
https://qemu-project.gitlab.io/qemu/system/images.html
- 30 <https://github.com/Metamogul/UTM-Snapshot-Manager>
https://github.com/Metamogul/UTM-Snapshot-Manager